

Python Workflow

Most of this document uses Patrick Pineault's 'The Good Research Code Handbook' . The sections about using NYU's HPC Greene cluster are from: 'Logging into HPC Clusters within NYU Network on Mac' . Transferring files to HPC Clusters from a Mac/Linux . Slurm Job Scheduler .

What is missing from this document: running JAX on the cluster.

Section A: Making the directory

1. Create the local directory called 'mydirectory'

```
mkdir mydirectory
```

2. Link the directory to Github inside VSCode

3. Create a Python virtual environment using conda

```
conda create -n envname python=3.11.4
conda activate envname
pip install jaxlib
pip install jax
conda list
conda info --envs
```

4. Initialise the environment in VSCode

Inside VSCode, click on 'Python' on the bottom right. Select 'Enter interpreter path'. Use the path from the `conda info --envs` step. Test the Python interpreter using a dummy `file.py` execution.

5. Export the environment; load the environment

Export:

```
conda env export > environment.yml
```

Recover:

```
conda env --name recoveredenv --file environment.yml
```

6. Make the project skeleton

```
mkdir {data,docs,results,scripts,src,tests}
```

- **data:** Where you put raw data for your project. You usually won't sync this to source control, unless you use very small, text-based datasets (< 10 MBs).

- **docs:** Where you put documentation, including Markdown and reStructuredText (reST).

Calling it docs makes it easy to publish documentation online through Github pages.

- **results:** Where you put results, including checkpoints, hdf5 files, pickle files, as well as figures and tables. If these files are heavy, you won't put these under source control.

- **scripts:** Where you put scripts - Python and bash alike - as well as .ipynb notebooks.

- **src:** Where you put reusable Python modules for your project. This is the kind of python code that you import.

- **tests:** Where you put tests for your code. We'll cover testing in a later lesson.

Section B: Saving and loading Python scripts

The project directory should include a few more files

```
mydirectory
- src
  - __init__.py
  - quantum.py
- setup.py
```

We can create these files from the command line

1. Initilisation file

```
touch __init__.py
```

2. Setup file

```
touch setup.py
nano setup.py
```

```
from setuptools import find_packages, setup
setup(
    name="src",
    packages=find_packages()
)
```

3. Run the packages

```
pip install -e .
```

4. Test the packages

```
python
import src.quantum
src.quantum
exit()
```

```
from src.quantum import *
print(sigma)
```

Section C: Running programs on NYU's Cluster

In this section, we run a Python file called `file.py` by executing the SBATCH file `run.s`

1. Creating the SBATCH file

```
touch run.s
```

In VSCode:

```
#!/bin/bash
#SBATCH --nodes=1
...
#SBATCH --output=slurm

module purge
module load python/intel/3.11.4
python ./file.py
```

I should set the exact SBATCH parameters based on my needs (GPU, RAM, CPU). The file `file.py` is what I want the cluster to run. I need to purge the node before loading python to ensure that I run the correct version.

2. Copy files to the cluster

```
cd dev/mydirectory
scp run.s file.py vv2102@greene.hpc.nyu.edu:~
... Enter password
ssh-keygen vv2102@greene.hpc.nyu.edu
ssh vv2102@greene.hpc.nyu.edu
... Enter password
mkdir test
mv run.s test/run.s
mv file.py test/file.py
```

To copy an entire folder, we use a similar command

```
scp -r mydirectory vv2102@greene.hpc.nyu.edu:~
```

3. Run the SBATCH file

```
cd test
sbatch run.sP
```

Section D: Running JAX on the Greene Cluster

This section follows NYU's guide, which you can find [here](#).

Running Python packages that are non-native to the Greene cluster, such as JAX, requires Singularity. In this document, I install JAX on the Greene cluster.

1. Create the project's directory in /scratch/...

Login to the cluster using the usual steps.

```
% ssh vv2102@greene.hpc.nyu.edu
# enter password
```

Now make the directory, which we call 'jax-example'. Note that my username is 'vv2102'.

```
$ mkdir /scratch/vv2102/jax-example
```

2. Copy an appropriate gzipped overlay image

```
$ cd /scratch/vv2102/jax-example
$ cp -rp /scratch/work/public/overlay-fs-ext3/overlay-15GB-500K.ext3.gz .
$ gunzip overlay-15GB-500K.ext3.gz
```

This copies a 15 GB and 500 000 file capacity overlay 'overlay-15GB-500K.ext3.gz' into the 'jax-example' folder. Overlay images allocate disk space.

We could have used a different overlay, which we can find inside the cluster's directory:

```
$ ls /scratch/work/public/overlay-fs-ext3
```

3. Load a Singularity container into the image

We load CUDA 12 on Ubuntu (cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.sif):

```
$ singularity exec --overlay overlay-15GB-500K.ext3:rw /scratch/work/public/
singularity/cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.sif /bin/bash
```

This line runs the Singularity application, opens the 15 GB .ext3 overlay image we created, loads the CUDA 12 .sif container into the image, and saves the .ext3 image.

We could have used a different Singularity container, which we can find inside the cluster's directory:

```
$ ls /scratch/work/public/singularity/
```

3. Installing miniconda inside Singularity

```
Singularity> wget https://repo.continuum.io/miniconda/Miniconda3-latest-Linux-x86_64.sh
# downloads files
Singularity> bash Miniconda3-latest-Linux-x86_64.sh -b -p /ext3/miniconda3
# downloads files
Singularity> rm Miniconda3-latest-Linux-x86_64.sh
# Remove unnecessary file
```

4. Create the conda wrapper

We place the wrapper insider /ext3/env.sh. The file does not exist yet, so we must create it.

```
Singularity> touch /ext3/env.sh
Singularity> nano /ext3/env.sh
```

The file contents of /ext3/env.sh are

```
#!/bin/bash

source /ext3/miniconda3/etc/profile.d/conda.sh
export PATH=/ext3/miniconda3/bin:$PATH
export PYTHONPATH=/ext3/miniconda3/bin:$PATH
```

Finally, we install all the relevant conda packages:

```
Singularity> source /ext3/env.sh
# activates conda
Singularity> conda update -n base conda -y
Singularity> conda clean --all --yes
Singularity> conda install pip -y
Singularity> conda install ipykernel -y # Note: ipykernel is required to run as a
kernel in the Open OnDemand Jupyter Notebooks
```

We can verify that everything is up to date

```
Singularity> conda --version
#conda
23.7.2
Singularity> python --version
#Python
3.11.4
Singularity> pip --version
#pip 23.2.1 from /ext3/miniconda3/lib/python3.11/site-packages/pip (python 3.11)
```

Exit Singularity.

```
Singularity> exit
```

5. Installing JAX (or other packages)

This is the only step not provided in NYU's guide. It is critical that JAX and the GPU have the same CUDA version. In my example, my GPU is CUDA 12, which was indicated in the Singularity container we loaded (`cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.sif`), so I shall install the CUDA 12 version of JAX.

Official JAX installation instructions are here.

First, I run my image on the cluster:

```
$ srun --cpus-per-task=2 --mem=10GB --time=04:00:00 --pty /bin/bash

# srun: job 36535589 queued and waiting for resources
# wait to be assigned a node
# srun: job 36535589 has been allocated resources

$ singularity exec --overlay overlay-15GB-500K.ext3:rw /scratch/work/public/
singularity/cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.sif /bin/bash

Singularity> source /ext3/env.sh
# activate the environment
```

Second, I install JAX (and a few other packages) using pip :

```
Singularity> pip install --upgrade pip
# Requirement already satisfied: pip in /ext3/miniconda3/lib/python3.11/site-
packages (23.2.1)

# CUDA 12 installation
# Note: wheels only available on linux.
Singularity> pip install --upgrade "jax[cuda12_pip]" -f https://
storage.googleapis.com/jax-releases/jax_cuda_releases.html

# Other packages
Singularity> pip3 install jupyter jupyterhub pandas matplotlib
```

We can verify that there is still room on our image:

```
Singularity> find /ext3 | wc -l
# 59867
Singularity> du -sh /ext3
# 4.1G /ext3
```

Lastly, I exit the Singularity container and rename the `.ext3` image into something more appropriate:

```
Singularity> exit
$ mv overlay-15GB-500K.ext3 my_jax.ext3
```

The `my_jax.ext3` image contains everything to run JAX on the cluster. We can now access Singularity and activate the environment using

```
$ singularity exec --overlay my_jax.ext3:rw /scratch/work/public/singularity/
cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.sif /bin/bash
Singularity> source /ext3/env.sh
Singularity> conda activate recoveredenv #running the virtual environment
recoveredenv
```

6. Running JAX

To run a Python script on the cluster, I include the necessary `#!/bin/env python` line of code called a shebang.

In this document, I run the `jax_detect_GPU.py` file as a test:

```
#!/bin/env python

from jax.lib import xla_bridge

if __name__ == "__main__":
    xla_backend = xla_bridge.get_backend()
    xla_backend_type = xla_bridge.get_backend().platform # cpu, gpu, tpu
    print(f"XLA backend type: {xla_backend_type}")

    gpu_count = xla_backend.device_count() if xla_backend_type == "gpu" else 0
    print(f"\nNumber of GPUs found on system: {gpu_count}")
    if xla_backend_type == "gpu":
        for idx, device in enumerate(xla_backend.devices()):
            gpu_type = "Active GPU" if idx == 0 else "GPU"
            print(f"\n{gpu_type} index: {device.id}")
            print(f"{gpu_type} name: {device.device_kind}")
```

I want to run `jax_detect_GPU.py` but I need to use a SLURM script. The SLURM script called `run-test.SBATCH` has contents:

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --time=1:00:00
#SBATCH --mem=2GB
#SBATCH --gres=gpu
#SBATCH --job-name=jax

module purge

singularity exec --nv \
    --overlay /scratch/vv2102/jax-example/my_jax.ext3:ro \
    /scratch/work/public/singularity/cuda12.1.1-cudnn8.9.0-devel-ubuntu22.04.2.
sif\
    /bin/bash -c "source /ext3/env.sh; python jax_detect_GPU.py"
```

Finally, I execute the script:


```
$ sbatch run-test.SBATCH
```

I can verify that JAX has worked by inspecting the file contents of the `slurm-xxxxxxx.out` file:

```
$ nano slurm-36537735.out

# XLA backend type: gpu
# Number of GPUs found on system: 1
# Active GPU index: 0
# Active GPU name: Tesla V100-PCIE-32GB
```

The file confirms JAX is running on GPU Tesla V100-PCIE-32GB.

Section E: Running custom environments on the Greene cluster

Julia on Greene Cluster

Create a directory for your julia work, such as /scratch//julia, and then change to your home directory. An example is shown below.

```
mkdir /home/<NetID>/julia  
cd /home/<NetID>/julia'
```

Copy an overlay image, such as the 2GB 100K overlay, which generally has enough storage for Julia packages. Once copied, unzip to the same folder, rename to julia-pkgs.ext3

```
cp -rp /scratch/work/public/overlay-fs-ext3/overlay-2GB-100K.ext3.gz .  
gunzip overlay-2GB-100K.ext3.gz  
mv overlay-2GB-100K.ext3 julia-pkgs.ext3
```

Copy the following wrapper script in the directory

```
cp -rp /share/apps/utis/julia-setup/* .
```

Now launch writable Singularity overlay, run the Julia REPL, and install packages

```
srun --cpus-per-task=2 --mem=10GB --time=04:00:00 --pty /bin/bash  
module purge  
bash my-julia-writable
```